

Contents

1	Motivation	4
1.1	The multi-agent gap in Hermes today	4
1.2	What others have built, and what broke	4
1.3	Google Gemini Enterprise: control vs. execution	5
1.4	Design thesis	5
2	Related Systems: Comparative Analysis	7
3	Architecture	8
3.1	Overview	8
3.2	Three-plane separation	8
3.3	The critical invariant: no in-process subagent swarms	8
4	Data Model	9
4.1	SQLite schema	9
4.2	Status semantics	10
4.3	Workspace kinds	10
5	Collaboration Patterns	11
5.1	P1 — Fan-out	11
5.2	P2 — Pipeline	11
5.3	P3 — Voting / Quorum (fan-in)	11
5.4	P4 — Long-running journal	11
5.5	P5 — Human-in-the-loop triage	11
5.6	P6 — @mention delegation	11
5.7	P7 — Thread-scoped workspace	11
5.8	P8 — Fleet farming	12
6	The Orchestrator Profile	13
6.1	Three properties of a well-behaved orchestrator	13
6.2	Shipping implication: installable profile templates	13
6.3	Why this is not kernel work	14
7	Multi-Tenant Context	15
7.1	The tenant primitive	15
7.2	What “tenant” actually means	15
7.3	What stays shared across tenants	15
7.4	What is deliberately <i>not</i> here	15
7.5	The “small company serving multiple clients” ask, solved	16
8	Worked Example: 50-Account Social Media Fleet	17
8.1	Setup	17
8.2	Per-tick task generation	17
8.3	The dispatcher does what it does	17
8.4	Why this is clean	18
9	User Stories	19
9.1	Story 1: Research triage and synthesis (non-coding)	19
9.2	Story 2: Scheduled recurring workflow (non-coding)	20
9.3	Story 3: Digital-twin / persistent assistant role (non-coding)	20
9.4	Story 4: Coding pipeline (code-shaped)	20